



Architecture

dual-head dilated CNN + Transformer · 5,071,106 parameters · PyTorch



one residual dilated block

identical for all 10 blocks; only the dilation d changes



objective

multinomial NLL over the profile

+ 0.1 x MSE on log total count

The profile head is scale-free: it predicts only the SHAPE of translation. Sequencing depth is isolated in the count head, so the transferable signal is not entangled with dataset-specific magnitude.