




a**Input: one transcript, per nucleotide (L up to 10,000 nt)****b**

5,071,106 trainable parameters | PyTorch | loss = multinomial NLL(profile) + 0.1 x MSE(log count)
28 epochs, 14.2 h on 1x A5500; early stop (patience 6) on median per-transcript Pearson; best epoch 21

c**d**