

Fig 2b -- the model selects a denser search space than naive enumeration, reproducibly across 5 datasets and both released models

