



a

Fig 2b -- the model selects a denser search space than naive enumeration, reproducibly across 5 datasets and both released models

**b**

D12 -- model-selected ORFs vs CPAT / CPC2 coding potential (released mamba4, frozen search params)

**c**

Predicting which ORFs to include keeps the DB compact: 2-3.4x higher discovery efficiency, far less CDS displacement, ~1/3 the decoy load vs the naive 3-frame null

